Difference and paired 95% bootstrap interval for JEV and Clef-Flash on the test images.
Decision model (k = 2 references per class; final description, pans: generic) minus the
CLIP baseline (k = 50; pans: k = 10); both use Platt scaling on the calibration set C.

